

B

bge-m3

Text Embeddings • baai

@cf/baai/bge-m3

Multi-Functionality, Multi-Linguality, and Multi-Granularity embeddings model.

Model Info

Batch	Yes
Unit Pricing	\$0.012 per M input tokens

Usage

- Workers - TypeScript

TypeScript

```
export interface Env {  
  AI: Ai;  
}  
  
export default {  
  async fetch(request, env):  
    Promise<Response> {  
  
    // Can be a string or array of strings  
    const stories = [  
      "This is a story about an orange  
cloud",  
      "This is a story about a llama",  
      "This is a story about a hugging  
emoji",  
    ];  
  
    const embeddings = await env.AI.run(  
      "@cf/baai/bge-m3",  
      {  
        text: stories,  
      }  
    );
```

```
    return Response.json(embeddings);  
  },  
} satisfies ExportedHandler<Env>;
```

▼ Python

Python

```
import os
import requests

ACCOUNT_ID = "your-account-id"
AUTH_TOKEN = os.environ.get("CLOUDFLARE_AUTH_TOKEN")

stories = [
    'This is a story about an orange cloud',
    'This is a story about a llama',
    'This is a story about a hugging emoji'
]

response = requests.post(
    f"https://api.cloudflare.com/client/v4/accounts/{ACCOUNT_ID}/workers-ai/bge-m3",
    headers={"Authorization": f"Bearer {AUTH_TOKEN}"},
    json={"text": stories}
)

print(response.json())
```

▼ curl

```
curl
https://api.cloudflare.com/client/v4/accounts/
m3 \
-X POST \
-H "Authorization: Bearer $CLOUDFLARE_API_TOKEN" \
-d '{ "text": ["This is a story about an orange",
"This is a story about a hugging emoji"] }'
```

① OpenAI compatible endpoints

Workers AI also supports OpenAI compatible API endpoints for `/v1/chat/completions` and `/v1/embeddings`. For more details, refer to [Configurations](#).

Parameters

* indicates a required field

Input

0 object

- query **string** min 1

A query you wish to perform against the provided contexts. If no query is provided the model will respond with embeddings for contexts

- contexts **array** required

List of provided contexts. Note that the index in this array is important, as the response will refer to it.

■ items **object**

- text **string** min 1

One of the provided context content

- truncate_inputs **boolean**

When provided with too long context should the model error out or truncate the context to fit?

1 object

- **text** **one of** required

- 0 **string** min 1

The text to embed

- 1 **array**

Batch of text values to embed

- **items** **string** min 1

The text to embed

- **truncate_inputs** **boolean**

When provided with too long context should the model error out or truncate the context to fit?

- 2 **object**

- **requests** **array** required

Batch of the embeddings requests to run using async-queue

- **items** **one of**

- 0 **object**

- **query** **string** min 1

A query you wish to perform against the provided contexts. If no query is provided the model will respond with embeddings for contexts

- **contexts** **array** required

List of provided contexts. Note that the index in this array is important, as the response will refer to it.

- **items** **object**

- **text** **string** min 1

One of the provided context content

- **truncate_inputs** **boolean**

When provided with too long context should the model error out or truncate the context to fit?

- **1** **object**

- **text** **one of** required

- **0** **string** min 1

The text to embed

- **1** **array**

Batch of text values to embed

- **items** **string** min 1

The text to embed

- **truncate_inputs** **boolean**

When provided with too long context should the model error out or truncate the context to fit?

Output

- **0** **object**

- **response** **array**

- **items** **object**

- **id** **integer**

Index of the context in the request

- **score** **number**

Score of the context under the index.

- ▶ **1** **object**

- **response** **array**

- **items** **array**

- **items** **number**

- **shape** **array**

- **items** **number**

- **pooling** **string**

The pooling method used in the embedding process.

- ▶ **2** **object**

- **shape** **array**

- **items** **number**

- **data** **array**

Embeddings of the requested text values

- **items** **array**

Floating point embedding representation
shaped by the embedding model

- **items** **number**

- **pooling** **string**

The pooling method used in the embedding
process.

- ▶ **3** **object**

- **request_id** **string**

The async request id that can be used to obtain the
results.

API Schemas

The following schemas are based on JSON Schema

Input

Output

```
{
  "type": "object",
  "oneOf": [
    {
      "title": "BGE M3 Input Query and
Contexts",
      "properties": {
        "query": {
          "type": "string",
          "minLength": 1,
          "description": "A query
you wish to perform against the provided
contexts. If no query is provided the model
with respond with embeddings for contexts"
        },
        "contexts": {
          "type": "array",
          "items": {
            "type": "object",
            "properties": {
              "text": {
                "type":
"string",
```

```
        "minLength":  
1,  
        "description":  
"One of the provided context content"  
    }  
}  
},  
    "description": "List of  
provided contexts. Note that the index in this  
array is important, as the response will refer  
to it."  
},  
    "truncate_inputs": {  
        "type": "boolean",  
        "default": false,  
        "description": "When  
provided with too long context should the  
model error out or truncate the context to  
fit?"  
    }  
},  
    "required": [  
        "contexts"  
    ]  
},
```

```
{
  "title": "BGE M3 Input Embedding",
  "properties": {
    "text": {
      "oneOf": [
        {
          "type": "string",
          "description":
            "The text to embed",
          "minLength": 1
        },
        {
          "type": "array",
          "description":
            "Batch of text values to embed",
          "items": {
            "type":
              "string",
            "description":
              "The text to embed",
            "minLength": 1
          },
          "maxItems": 100
        }
      ]
    }
  }
}
```



```
    },  
    "truncate_inputs": {  
      "type": "boolean",  
      "default": false,  
      "description": "When  
provided with too long context should the  
model error out or truncate the context to  
fit?"  
    }  
  },  
  "required": [  
    "text"  
  ]  
},  
{  
  "properties": {  
    "requests": {  
      "type": "array",  
      "description": "Batch of  
the embeddings requests to run using async-  
queue",  
      "items": {  
        "type": "object",  
        "oneOf": [  
          {
```

```
      "title": "BGE  
M3 Input Query and Contexts",  
      "properties":  
      {  
        "query": {  
  
          "type": "string",  
  
          "minLength": 1,  
  
          "description": "A query you wish to perform  
against the provided contexts. If no query is  
provided the model will respond with  
embeddings for contexts"  
        },  
  
        "contexts": {  
  
          "type": "array",  
  
          "items": {  
  
            "type": "object",  
  
            "properties": {
```

```
"text": {  
  
  "type": "string",  
  
  "minLength": 1,  
  
  "description": "One of the provided context  
content"  
  
}  
  
},  
  
"description": "List of provided contexts.  
Note that the index in this array is  
important, as the response will refer to it."  
},  
  
"truncate_inputs": {  
  
  "type": "boolean",  
  
  "default": false,
```

```
"description": "When provided with too long  
context should the model error out or truncate  
the context to fit?"
```

```
    }  
  },  
  "required": [  
    "contexts"  
  ],  
},  
{  
  "title": "BGE  
M3 Input Embedding",  
  "properties":  
{  
    "text": {  
      "oneOf": [  
        {  
          "type": "string",  
          "description": "The text to embed",  
          "minLength": 1
```

```
    },  
    {  
  
      "type": "array",  
  
      "description": "Batch of text values to  
embed",  
  
      "items": {  
  
        "type": "string",  
  
        "description": "The text to embed",  
  
        "minLength": 1  
  
      },  
  
      "maxItems": 100  
  
    }  
  ]  
},  
  
  "truncate_inputs": {
```

```
"type": "boolean",

"default": false,

"description": "When provided with too long
context should the model error out or truncate
the context to fit?"

    }
  },
  "required": [
    "text"
  ]
}

]
}

},
"required": [
  "requests"
]
}

]
}
```

Was this helpful?



Previous

← [Dashboard](#)

Next

[Workers Bindings](#) →

Resources

[API](#)

[New to Cloudflare?](#)

[Directory](#)

[Sponsorships](#)

[Open Source](#)

Support

[Help Center](#)

[System Status](#)

[Compliance](#)

[GDPR](#)

Company

[cloudflare.com](#)

[Our team](#)

[Careers](#)

Tools

[Cloudflare Radar](#)

[Speed Test](#)

[Is BGP Safe Yet?](#)

[RPKI Toolkit](#)

Certificate Transparency

Community



© 2025 Cloudflare, Inc. • [Privacy Policy](#) • [Terms of Use](#)

• [Report Security Issues](#) • [Trademark](#) •  [Cookie Preferences](#)