



B

bge-small-en-v1.5

Text Embeddings • baai

@cf/baai/bge-small-en-v1.5

BAAI general embedding (Small) model that transforms any given text into a 384-dimensional vector

Model Info

More information

[link ↗](#)

Maximum Input
Tokens

512

Output Dimensions

384

Batch

Yes

Model Info

Unit Pricing

\$0.02 per M input
tokens

Usage

✓ Workers - TypeScript

TypeScript

```
export interface Env {  
  AI: Ai;  
}  
  
export default {  
  async fetch(request, env):  
    Promise<Response> {  
  
    // Can be a string or array of strings  
    const stories = [  
      "This is a story about an orange  
cloud",  
      "This is a story about a llama",  
      "This is a story about a hugging  
emoji",  
    ];  
  
    const embeddings = await env.AI.run(  
      "@cf/baai/bge-small-en-v1.5",  
      {  
        text: stories,  
      }  
    );  
  }  
};
```

```
    return Response.json(embeddings);  
  },  
} satisfies ExportedHandler<Env>;
```

▼ Python

Python

```
import os
import requests

ACCOUNT_ID = "your-account-id"
AUTH_TOKEN = os.environ.get("CLOUDFLARE_AUTH_TOKEN")

stories = [
    'This is a story about an orange cloud',
    'This is a story about a llama',
    'This is a story about a hugging emoji'
]

response = requests.post(
    f"https://api.cloudflare.com/client/v4/accounts/{ACCOUNT_ID}/workers-ai/bge-small-en-v1.5",
    headers={"Authorization": f"Bearer {AUTH_TOKEN}"},
    json={"text": stories}
)

print(response.json())
```

▼ curl

```
curl
```

```
https://api.cloudflare.com/client/v4/accounts/  
small-en-v1.5 \\\n-X POST \\\n-H "Authorization: Bearer $CLOUDFLARE_API_TOKEN" \\\n-d '{ "text": ["This is a story about an orange",  
"This is a story about a hugging emoji"] }'
```

① OpenAI compatible endpoints

Workers AI also supports OpenAI compatible API endpoints for `/v1/chat/completions` and `/v1/embeddings`. For more details, refer to [Configurations](#).

Parameters

* indicates a required field

Input

- 0 **object**

- text **one of** required

- 0 **string** min 1

The text to embed

- 1 **array**

Batch of text values to embed

- items **string** min 1

The text to embed

- pooling **string** default mean

The pooling method used in the embedding process. `cls` pooling will generate more accurate embeddings on larger inputs - however, embeddings created with cls pooling are not compatible with embeddings generated with mean pooling. The default pooling method is `mean` in order for this to not be a breaking change, but we

highly suggest using the new `cls` pooling for better accuracy.

► 1 **object**

- requests **array** required

Batch of the embeddings requests to run using `async-queue`

■ items **object**

- text **one of** required

- 0 **string** min 1

The text to embed

- 1 **array**

Batch of text values to embed

- items **string** min 1

The text to embed

- pooling **string** default mean

The pooling method used in the embedding process. `cls` pooling will generate more

accurate embeddings on larger inputs - however, embeddings created with `cls` pooling are not compatible with embeddings generated with mean pooling. The default pooling method is ``mean`` in order for this to not be a breaking change, but we highly suggest using the new ``cls`` pooling for better accuracy.

Output

- `0` **object**

- `shape` **array**

- `items` **number**

- `data` **array**

Embeddings of the requested text values

- `items` **array**

Floating point embedding representation
shaped by the embedding model

- items **number**

- pooling **string**

The pooling method used in the embedding process.

- ▶ 1 **object**

- request_id **string**

The async request id that can be used to obtain the results.

API Schemas

The following schemas are based on JSON Schema

Input

Output

```
{
  "type": "object",
  "oneOf": [
    {
      "properties": {
        "text": {
          "oneOf": [
            {
              "type": "string",
              "description":
                "The text to embed",
              "minLength": 1
            },
            {
              "type": "array",
              "description":
                "Batch of text values to embed",
              "items": {
                "type":
                  "string",
                "description":
                  "The text to embed",
                "minLength": 1
              },
            }
          ]
        }
      }
    }
  ]
}
```

```
        "maxItems": 100
      }
    ]
  },
  "pooling": {
    "type": "string",
    "enum": [
      "mean",
      "cls"
    ],
    "default": "mean",
    "description": "The
pooling method used in the embedding process.
`cls` pooling will generate more accurate
embeddings on larger inputs - however,
embeddings created with cls pooling are not
compatible with embeddings generated with mean
pooling. The default pooling method is `mean`
in order for this to not be a breaking change,
but we highly suggest using the new `cls`
pooling for better accuracy."
  }
},
"required": [
  "text"
```

```
    ]  
  },  
  {  
    "properties": {  
      "requests": {  
        "type": "array",  
        "description": "Batch of  
the embeddings requests to run using async-  
queue",  
        "items": {  
          "properties": {  
            "text": {  
              "oneOf": [  
                {  
  
"type": "string",  
  
"description": "The text to embed",  
  
"minLength": 1  
                },  
                {  
  
"type": "array",
```

```
"description": "Batch of text values to
embed",

"items": {

"type": "string",

"description": "The text to embed",

"minLength": 1

},

"maxItems": 100

}

],

},

"pooling": {
  "type":
"string",

  "enum": [
    "mean",
    "cls"
  ],
```



```
        "default":  
        "mean",  
        "description":  
        "The pooling method used in the embedding  
        process. `cls` pooling will generate more  
        accurate embeddings on larger inputs -  
        however, embeddings created with cls pooling  
        are not compatible with embeddings generated  
        with mean pooling. The default pooling method  
        is `mean` in order for this to not be a  
        breaking change, but we highly suggest using  
        the new `cls` pooling for better accuracy."  
    },  
    "required": [  
        "text"  
    ]  
},  
},  
"required": [  
    "requests"  
]  
}
```



Was this helpful?



Previous

← Dashboard

Next

Workers Bindings →

Resources

[API](#)

[New to Cloudflare?](#)

[Directory](#)

[Sponsorships](#)

[Open Source](#)

Support

[Help Center](#)

[System Status](#)

[Compliance](#)

[GDPR](#)

Company

[cloudflare.com](#)

[Our team](#)

Tools

[Cloudflare Radar](#)

[Speed Test](#)

Careers

Is BGP Safe Yet?

RPKI Toolkit

Certificate Transparency

Community

 X

 Discord

 YouTube

 GitHub

© 2025 Cloudflare, Inc. • [Privacy Policy](#) • [Terms of Use](#)

• [Report Security Issues](#) • [Trademark](#) •  [Cookie Preferences](#)