



A

gemma-sea-lion-v4-27b-it

Text Generation • aisingapore

@cf/aisingapore/gemma-sea-lion-v4-27b-it

SEA-LION stands for Southeast Asian Languages In One Network, which is a collection of Large Language Models (LLMs) which have been pretrained and instruct-tuned for the Southeast Asia (SEA) region.

Model Info

Context 128,000 tokens

Window [↗](#)

Unit Pricing	\$0.35 per M input tokens, \$0.56 per M output tokens
--------------	--

Playground

Try out this model with Workers AI LLM Playground. It does not require any setup or authentication and an instant way to preview and test a model directly in the browser.

[Launch the LLM Playground](#)

Usage

▼ Worker - Streaming

TypeScript

```
export interface Env {  
  AI: Ai;  
}  
  
export default {  
  async fetch(request, env):  
    Promise<Response> {  
  
    const messages = [  
      { role: "system", content: "You are a  
friendly assistant" },  
      {  
        role: "user",  
        content: "What is the origin of the  
phrase Hello, World",  
      },  
    ];  
  
    const stream = await  
env.AI.run("@cf/aisingapore/gemma-sea-lion-  
v4-27b-it", {  
      messages,  
      stream: true,  
    });
```

```
return new Response(stream, {  
  headers: { "content-type":  
"text/event-stream" },  
  });  
},  
} satisfies ExportedHandler<Env>;
```

✓ Worker

TypeScript

```
export interface Env {  
  AI: Ai;  
}  
  
export default {  
  async fetch(request, env):  
    Promise<Response> {  
  
    const messages = [  
      { role: "system", content: "You are a  
friendly assistant" },  
      {  
        role: "user",  
        content: "What is the origin of the  
phrase Hello, World",  
      },  
    ];  
  
    const response = await  
env.AI.run("@cf/aisingapore/gemma-sea-lion-  
v4-27b-it", { messages });  
  
    return Response.json(response);  
  },  
} satisfies ExportedHandler<Env>;
```

▼ Python

Python

```
import os
import requests

ACCOUNT_ID = "your-account-id"
AUTH_TOKEN = os.environ.get("CLOUDFLARE_AUTH_TOKEN")

prompt = "Tell me all about PEP-8"
response = requests.post(
    f"https://api.cloudflare.com/client/v4/accounts/{ACCOUNT_ID}/workers-ai/generate",
    headers={"Authorization": f"Bearer {AUTH_TOKEN}"},
    json={
        "messages": [
            {"role": "system", "content": "You are a helpful assistant."},
            {"role": "user", "content": prompt}
        ]
    }
)
result = response.json()
print(result)
```

▼ curl

```
curl
```

```
https://api.cloudflare.com/client/v4/accounts/  
sea-lion-v4-27b-it \  
-X POST \  
-H "Authorization: Bearer $CLOUDFLARE_AUTH_TOKEN" \  
-d '{ "messages": [{ "role": "system", "content": "You are a helpful assistant." },  
"user", "content": "Why is pizza so good" } ] }'
```

① OpenAI compatible endpoints

Workers AI also supports OpenAI compatible API endpoints for `/v1/chat/completions` and `/v1/embeddings`. For more details, refer to [Configurations](#).

Parameters

* indicates a required field

Input

- ▶ 0 **object**
 - \$merge
- ▶ 1 **object**
 - \$merge
- ▶ 2 **object**
 - requests **array** required
 - items **one of**
 - 0 **object**
 - \$merge
 - 1 **object**
 - \$merge

Output

- ▶ **0** **object**
 - **id** **string**

Unique identifier for the completion
 - **object** **string**

Object type identifier
 - **created** **number**

Unix timestamp of when the completion was created
 - **model** **string**

Model used for the completion
 - **choices** **array**

List of completion choices
 - **items** **object**
 - **index** **number**

Index of the choice in the list

- **message** **object**

The message generated by the model

- **role** **string** required

Role of the message author

- **content** **string** required

The content of the message

- **reasoning_content** **string**

Internal reasoning content (if available)

- **tool_calls** **array**

Tool calls made by the assistant

- **items** **object**

- **id** **string** required

Unique identifier for the tool call

- **type** **string** required

Type of tool call

- **function** **object** required

- **name** **string** required

Name of the function to call

- **arguments** **string** required

JSON string of arguments for the function

- **finish_reason** **string**

Reason why the model stopped generating

- **stop_reason** **string**

Stop reason (may be null)

- **logprobs** **object**

Log probabilities (if requested)

- **usage**

- **prompt_logprobs** **object**

Log probabilities for the prompt (if requested)

- **1** **object**

- **id** **string**

Unique identifier for the completion

- **object** **string**

Object type identifier

- **created** **number**

Unix timestamp of when the completion was created

- **model** **string**

Model used for the completion

- **choices** **array**

List of completion choices

- **items** **object**

- **index** **number** required

Index of the choice in the list

- **text** **string** required

The generated text completion

- **finish_reason** **string** required

Reason why the model stopped generating

- `stop_reason` **string**

Stop reason (may be null)

- `logprobs` **object**

Log probabilities (if requested)

- `prompt_logprobs` **object**

Log probabilities for the prompt (if requested)

- `usage`

- 2 **string**

- 3 **object**

- `request_id` **string**

The async request id that can be used to obtain the results.

API Schemas

The following schemas are based on JSON Schema

Input

Output

```
{
  "$id":
  "http://ai.cloudflare.com/schemas/textGenerationI
  "type": "object",
  "oneOf": [
    {
      "title": "Prompt",
      "properties": {
        "$merge": {
          "source": {
            "prompt": {
              "$ref":
                "textGenerationPrompts#/prompt"
            },
            "lora": {
              "$ref":
                "textGenerationFinetune#/lora"
            },
            "response_format": {
              "$ref":
                "jsonMode#/response_format"
            }
          },
          "with": {
```

```
    "raw": {
      "type": "boolean",
      "default": false,
      "description": "If the chat template is not applied and you must adhere to the specific model's expected formatting."
    },
    "stream": {
      "type": "boolean",
      "default": false,
      "description": "If true, the response will be streamed back incrementally using SSE, Server Sent Events."
    },
    "max_tokens": {
      "type": "integer",
      "default": 2000,
      "description": "The maximum number of tokens to generate in the response."
    },
    "temperature": {
      "type": "number",
      "default": 0.6,
      "minimum": 0,
      "maximum": 5,
```

```
    "description": "Controls the randomness of the output; higher values produce more random results."
```

```
  },
```

```
  "top_p": {
```

```
    "type": "number",
```

```
    "minimum": 0.001,
```

```
    "maximum": 1,
```

```
    "description": "Adjusts the creativity of the AI's responses by controlling many possible words it considers. Lower values make outputs more predictable; higher values allow for varied and creative responses."
```

```
  },
```

```
  "top_k": {
```

```
    "type": "integer",
```

```
    "minimum": 1,
```

```
    "maximum": 50,
```

```
    "description": "Limits AI to choose from the top 'k' most probable words. Lower values make responses more focused; higher values introduce more variety and potential surprises."
```

```
  },
```

```
  "seed": {
```

```
    "type": "integer",
```

```
        "minimum": 1,  
        "maximum": 9999999999,  
        "description": "Random seed for reproducibility of the generation."  
    },  
    "repetition_penalty": {  
        "type": "number",  
        "minimum": 0,  
        "maximum": 2,  
        "description": "Penalty for repeated tokens; higher values discourage repetition."  
    },  
    "frequency_penalty": {  
        "type": "number",  
        "minimum": -2,  
        "maximum": 2,  
        "description": "Decrease the likelihood of the model repeating the same list of words verbatim."  
    },  
    "presence_penalty": {  
        "type": "number",  
        "minimum": -2,  
        "maximum": 2,
```

```
        "description": "Increase  
the likelihood of the model introducing new topics",  
      },  
    ],  
  },  
  "required": [  
    "prompt"  
  ],  
},  
{  
  "title": "Messages",  
  "properties": {  
    "$merge": {  
      "source": {  
        "messages": {  
          "$ref":  
"textGenerationPrompts#/messages"  
        },  
        "functions": {  
          "$ref":  
"textGenerationTools#/functions"  
        },  
        "tools": {
```

```
        "$ref":  
"textGenerationTools#/tools"  
    },  
    "response_format": {  
        "$ref":  
"jsonMode#/response_format"  
    },  
    },  
    "with": {  
        "raw": {  
            "type": "boolean",  
            "default": false,  
            "description": "If tr  
chat template is not applied and you must adhere  
specific model's expected formatting."  
        },  
        "stream": {  
            "type": "boolean",  
            "default": false,  
            "description": "If tr  
the response will be streamed back incrementally  
SSE, Server Sent Events."  
        },  
        "max_tokens": {  
            "type": "integer",
```

```
        "default": 2000,  
        "description": "The  
maximum number of tokens to generate in the response",  
    },  
    "temperature": {  
        "type": "number",  
        "default": 0.6,  
        "minimum": 0,  
        "maximum": 5,  
        "description": "Controls  
the randomness of the output; higher values produce  
more random results."  
    },  
    "top_p": {  
        "type": "number",  
        "minimum": 0.001,  
        "maximum": 1,  
        "description": "Adjusts  
the creativity of the AI's responses by controlling  
many possible words it considers. Lower values make  
outputs more predictable; higher values allow for  
varied and creative responses."  
    },  
    "top_k": {  
        "type": "integer",
```



```
        "minimum": 1,  
        "maximum": 50,  
        "description": "Limit  
AI to choose from the top 'k' most probable words  
Lower values make responses more focused; higher  
introduce more variety and potential surprises."  
    },  
    "seed": {  
        "type": "integer",  
        "minimum": 1,  
        "maximum": 9999999999,  
        "description": "Random  
seed for reproducibility of the generation."  
    },  
    "repetition_penalty": {  
        "type": "number",  
        "minimum": 0,  
        "maximum": 2,  
        "description": "Penal  
for repeated tokens; higher values discourage  
repetition."  
    },  
    "frequency_penalty": {  
        "type": "number",  
        "minimum": -2,
```

```
        "maximum": 2,
        "description": "Decrease
the likelihood of the model repeating the same li
verbatim."
    },
    "presence_penalty": {
        "type": "number",
        "minimum": -2,
        "maximum": 2,
        "description": "Increa
the likelihood of the model introducing new topic
    }
}
},
"required": [
    "messages"
]
},
{
    "title": "Async Batch",
    "type": "object",
    "properties": {
        "requests": {
            "type": "array",
```

```
      "items": {
        "type": "object",
        "oneOf": [
          {
            "title": "Prompt"
            "properties": {
              "$merge": {
                "source":
                  "prompt"
            }
          }
        ],
        "textGenerationPrompts#/prompt"
      },
      {
        "title": "Lora"
        "textGenerationFinetune#/lora"
      },
      {
        "title": "Response Format"
        "response_format": {
          "jsonMode#/response_format"
        }
      },
      {
        "title": "With"
        "with": {
```

"textGenerationOptions#/common"

"textGenerationPrompts#/messages"

"functions": {

"textGenerationTools#/functions"

```
"textGenerationTools#/tools"
```

```
},
```

```
"response_format": {
```

```
"jsonMode#/response_format"
```

```
}
```

```
},
```

```
"with": {
```

```
"$ref
```

```
"textGenerationOptions#/common"
```

```
}
```

```
}
```

```
},
```

```
"required": [
```

```
"messages"
```

```
]
```

```
}
```

```
]
```

```
}
```

```
}
```

```
},
```

```
"required": [
```

```
"requests"
```

```
}  
]  
}  
]  
]
```

Was this helpful?



Previous

← [Dashboard](#)

Next

[Workers Bindings](#) →

Resources

[API](#)

[New to Cloudflare?](#)

[Directory](#)

[Sponsorships](#)

[Open Source](#)

Support

[Help Center](#)

[System Status](#)

[Compliance](#)

[GDPR](#)

Company

Tools

[cloudflare.com](#)[Cloudflare Radar](#)[Our team](#)[Speed Test](#)[Careers](#)[Is BGP Safe Yet?](#)[RPKI Toolkit](#)[Certificate Transparency](#)

Community

[!\[\]\(51514032c8ca341817228f39f1307b05_img.jpg\) X](#)[!\[\]\(c444627dab9fee9a1550c053ffaaaae2_img.jpg\) Discord](#)[!\[\]\(0d7ca0919e6c47bbd874bfa0189fe22e_img.jpg\) YouTube](#)[!\[\]\(274fd520e03b61c1b9ffc861754cacdc_img.jpg\) GitHub](#)

© 2025 Cloudflare, Inc. • [Privacy Policy](#) • [Terms of Use](#)

• [Report Security Issues](#) • [Trademark](#) •  [Cookie Preferences](#)