



llama-3.2-11b-vision-instruct

Text Generation • Meta

@cf/meta/llama-3.2-11b-vision-instruct

The Llama 3.2-Vision instruction-tuned models are optimized for visual recognition, image reasoning, captioning, and answering general questions about an image.

Note

To use Llama 3.2 11b Vision Instruct, you need to agree to the [Meta License](#) and [Acceptable Use Policy](#). To do so, please send an initial request to @cf/meta/llama-3.2-11b-vision-instruct with "prompt" : "agree". After that, you'll be able to use the model as normal.

```
curl
https://api.cloudflare.com/client/v4/accounts/
3.2-11b-vision-instruct \
-X POST \
-H "Authorization: Bearer $CLOUDFLARE_API_KEY" \
-d '{ "prompt": "agree" }'
```

Model Info

Context 128,000 tokens

Window [↗](#)

Model Info

Terms and
License

[link ↗](#)

LoRA

Yes

Unit Pricing

\$0.049 per M input tokens,
\$0.68 per M output tokens

Playground

Try out this model with Workers AI LLM Playground. It does not require any setup or authentication and an instant way to preview and test a model directly in the browser.

Launch the LLM Playground

Usage

✓ Worker - Streaming

TypeScript

```
export interface Env {  
  AI: Ai;  
}  
  
export default {  
  async fetch(request, env):  
    Promise<Response> {  
  
    const messages = [  
      { role: "system", content: "You are a  
friendly assistant" },  
      {  
        role: "user",  
        content: "What is the origin of the  
phrase Hello, World",  
      },  
    ];  
  
    const stream = await  
env.AI.run("@cf/meta/llama-3.2-11b-vision-  
instruct", {  
      messages,  
      stream: true,  
    });
```

```
return new Response(stream, {  
  headers: { "content-type":  
"text/event-stream" },  
  });  
},  
} satisfies ExportedHandler<Env>;
```

✓ Worker

TypeScript


```
export interface Env {  
  AI: Ai;  
}  
  
export default {  
  async fetch(request, env):  
    Promise<Response> {  
  
    const messages = [  
      { role: "system", content: "You are a  
friendly assistant" },  
      {  
        role: "user",  
        content: "What is the origin of the  
phrase Hello, World",  
      },  
    ];  
  
    const response = await  
env.AI.run("@cf/meta/llama-3.2-11b-vision-  
instruct", { messages });  
  
    return Response.json(response);  
  },  
} satisfies ExportedHandler<Env>;
```

▼ Python

Python

```
import os
import requests

ACCOUNT_ID = "your-account-id"
AUTH_TOKEN = os.environ.get("CLOUDFLARE_AUTH_TOKEN")

prompt = "Tell me all about PEP-8"
response = requests.post(
    f"https://api.cloudflare.com/client/v4/accounts/{ACCOUNT_ID}/llama-3.2-11b-vision-instruct",
    headers={"Authorization": f"Bearer {AUTH_TOKEN}"},
    json={
        "messages": [
            {"role": "system", "content": "You are a helpful assistant."},
            {"role": "user", "content": prompt}
        ]
    }
)
result = response.json()
print(result)
```

▼ curl

```
curl
```

```
https://api.cloudflare.com/client/v4/accounts/  
3.2-11b-vision-instruct \  
-X POST \  
-H "Authorization: Bearer $CLOUDFLARE_AUTH_TOKEN" \  
-d '{ "messages": [{ "role": "system", "content": "You are a helpful assistant." },  
"user", "content": "Why is pizza so good" ]}]}'
```

① OpenAI compatible endpoints

Workers AI also supports OpenAI compatible API endpoints for `/v1/chat/completions` and `/v1/embeddings`. For more details, refer to [Configurations](#).

Parameters

* indicates a required field

Input

▸ 0 **object**

- **prompt** **string** required min 1 max 131072

The input text prompt for the model to generate a response.

- **image** **one of**

■ 0 **array**

An array of integers that represent the image data constrained to 8-bit unsigned integer values. Deprecated, use image as a part of messages now.

■ **items** **number**

A value between 0 and 255

■ 1 **string**

Binary string representing the image contents. Deprecated, use image as a part of messages now.

- `raw` **boolean**

If true, a chat template is not applied and you must adhere to the specific model's expected formatting.

- `stream` **boolean**

If true, the response will be streamed back incrementally using SSE, Server Sent Events.

- `max_tokens` **integer** default 256

The maximum number of tokens to generate in the response.

- `temperature` **number** default 0.6 min 0 max 5

Controls the randomness of the output; higher values produce more random results.

- `top_p` **number** min 0 max 2

Adjusts the creativity of the AI's responses by controlling how many possible words it considers. Lower values make outputs more predictable; higher values allow for more varied and creative responses.

- `top_k` **integer** min 1 max 50

Limits the AI to choose from the top 'k' most probable words. Lower values make responses more focused; higher values introduce more variety and potential surprises.

- `seed` **integer** min 1 max 9999999999

Random seed for reproducibility of the generation.

- `repetition_penalty` **number** min 0 max 2

Penalty for repeated tokens; higher values discourage repetition.

- `frequency_penalty` **number** min 0 max 2

Decreases the likelihood of the model repeating the same lines verbatim.

- `presence_penalty` **number** min 0 max 2

Increases the likelihood of the model introducing new topics.

- `lora` **string**

Name of the LoRA (Low-Rank Adaptation) model to fine-tune the base model.

1 **object**

○ **messages** **array** required

An array of message objects representing the conversation history.

■ **items** **object**

■ **role** **string**

The role of the message sender (e.g., 'user', 'assistant', 'system', 'tool').

■ **tool_call_id** **string**

The tool call id. Must be supplied for tool calls for Mistral-3. If you don't know what to put here you can fall back to 000000001

■ **content** **one of**

■ **0** **string**

The content of the message as a string.

- 1 **array**
 - **items** **object**
 - **type** **string**

Type of the content provided
 - **text** **string**
 - **image_url** **object**
 - **url** **string**

image uri with data (e.g.
...).
HTTP URL will not be accepted
- 2 **object**
 - **type** **string**

Type of the content provided
 - **text** **string**
 - **image_url** **object**

- **url** **string**

image uri with data (e.g.
...).
HTTP URL will not be accepted

- **image** **one of**

- **0** **array**

An array of integers that represent the image data constrained to 8-bit unsigned integer values. Deprecated, use image as a part of messages now.

- **items** **number**

A value between 0 and 255

- **1** **string**

Binary string representing the image contents. Deprecated, use image as a part of messages now.

- **functions** **array**

- **items** **object**

- **name** **string** required

- **code** **string** required

- **tools** **array**

A list of tools available for the assistant to use.

- **items** **one of**

- **0** **object**

- **name** **string** required

The name of the tool. More descriptive the better.

- **description** **string** required

A brief description of what the tool does.

- **parameters** **object** required

Schema defining the parameters accepted by the tool.

- **type** **string** required

The type of the parameters object (usually 'object').

- **required** **array**

List of required parameter names.

- **items** **string**

- **properties** **object** required

Definitions of each parameter.

- **additionalProperties** **object**

- **type** **string** required

The data type of the parameter.

- **description** **string** required

A description of the expected parameter.

- **1** **object**

- **type** **string** required

Specifies the type of tool (e.g., 'function').

- **function** **object** required

Details of the function tool.

- **name** **string** required
The name of the function.
- **description** **string** required
A brief description of what the function does.
- **parameters** **object** required
Schema defining the parameters accepted by the function.
 - **type** **string** required
The type of the parameters object (usually 'object').
 - **required** **array**
List of required parameter names.
 - **items** **string**
 - **properties** **object** required
Definitions of each parameter.
 - **additionalProperties** **object**

- **type** **string** required

The data type of the parameter.

- **description** **string** required

A description of the expected parameter.

- **stream** **boolean**

If true, the response will be streamed back incrementally.

- **max_tokens** **integer** default 256

The maximum number of tokens to generate in the response.

- **temperature** **number** default 0.6 min 0 max 5

Controls the randomness of the output; higher values produce more random results.

- **top_p** **number** min 0 max 2

Controls the creativity of the AI's responses by adjusting how many possible words it considers.

Lower values make outputs more predictable; higher values allow for more varied and creative responses.

- `top_k` **integer** min 1 max 50

Limits the AI to choose from the top 'k' most probable words. Lower values make responses more focused; higher values introduce more variety and potential surprises.

- `seed` **integer** min 1 max 9999999999

Random seed for reproducibility of the generation.

- `repetition_penalty` **number** min 0 max 2

Penalty for repeated tokens; higher values discourage repetition.

- `frequency_penalty` **number** min 0 max 2

Decreases the likelihood of the model repeating the same lines verbatim.

- `presence_penalty` **number** min 0 max 2

Increases the likelihood of the model introducing new topics.

Output

- **0** **object**

- **response** **string**

The generated text response from the model

- **tool_calls** **array**

An array of tool calls requests made during the response generation

- **items** **object**

- **arguments** **object**

The arguments passed to be passed to the tool call request

- **name** **string**

The name of the tool to be called

- **1** **string**

API Schemas

The following schemas are based on JSON Schema

Input

Output

```
{
  "type": "object",
  "oneOf": [
    {
      "title": "Prompt",
      "properties": {
        "prompt": {
          "type": "string",
          "minLength": 1,
          "maxLength": 131072,
          "description": "The input
text prompt for the model to generate a
response."
        },
        "image": {
          "oneOf": [
            {
              "type": "array",
              "description": "An
array of integers that represent the image
data constrained to 8-bit unsigned integer
values.  Deprecated, use image as a part of
messages now.",
              "items": {
```

```
        "type":  
        "number",  
        "description":  
        "A value between 0 and 255"  
    },  
    {  
        "type": "string",  
        "format":  
        "binary",  
        "description":  
        "Binary string representing the image  
        contents.  Deprecated, use image as a part of  
        messages now."  
    }  
],  
,  
"raw": {  
    "type": "boolean",  
    "default": false,  
    "description": "If true, a  
chat template is not applied and you must  
adhere to the specific model's expected  
formatting."  
},
```

```
"stream": {  
  "type": "boolean",  
  "default": false,  
  "description": "If true,  
the response will be streamed back  
incrementally using SSE, Server Sent Events."  
},  
"max_tokens": {  
  "type": "integer",  
  "default": 256,  
  "description": "The  
maximum number of tokens to generate in the  
response."  
},  
"temperature": {  
  "type": "number",  
  "default": 0.6,  
  "minimum": 0,  
  "maximum": 5,  
  "description": "Controls  
the randomness of the output; higher values  
produce more random results."  
},  
"top_p": {  
  "type": "number",
```

```
        "minimum": 0,  
        "maximum": 2,  
        "description": "Adjusts  
the creativity of the AI's responses by  
controlling how many possible words it  
considers. Lower values make outputs more  
predictable; higher values allow for more  
varied and creative responses."  
    },  
    "top_k": {  
        "type": "integer",  
        "minimum": 1,  
        "maximum": 50,  
        "description": "Limits the  
AI to choose from the top 'k' most probable  
words. Lower values make responses more  
focused; higher values introduce more variety  
and potential surprises."  
    },  
    "seed": {  
        "type": "integer",  
        "minimum": 1,  
        "maximum": 9999999999,  
        "description": "Random  
seed for reproducibility of the generation."
```

```
    },  
    "repetition_penalty": {  
      "type": "number",  
      "minimum": 0,  
      "maximum": 2,  
      "description": "Penalty  
for repeated tokens; higher values discourage  
repetition."  
    },  
    "frequency_penalty": {  
      "type": "number",  
      "minimum": 0,  
      "maximum": 2,  
      "description": "Decreases  
the likelihood of the model repeating the same  
lines verbatim."  
    },  
    "presence_penalty": {  
      "type": "number",  
      "minimum": 0,  
      "maximum": 2,  
      "description": "Increases  
the likelihood of the model introducing new  
topics."  
    },  
  },
```

```
      "lora": {
        "type": "string",
        "description": "Name of
the LoRA (Low-Rank Adaptation) model to fine-
tune the base model."
      },
      "required": [
        "prompt"
      ],
    },
    {
      "title": "Messages",
      "properties": {
        "messages": {
          "type": "array",
          "description": "An array
of message objects representing the
conversation history.",
          "items": {
            "type": "object",
            "properties": {
              "role": {
                "type":
"string",
```



```
        "description":  
        "The role of the message sender (e.g., 'user',  
        'assistant', 'system', 'tool')."  
    },  
    "tool_call_id": {  
        "type":  
        "string",  
        "description":  
        "The tool call id. Must be supplied for tool  
        calls for Mistral-3. If you don't know what to  
        put here you can fall back to 000000001",  
        "pattern": "  
[a-zA-Z0-9]{9}"  
    },  
    "content": {  
        "oneOf": [  
        {  
        "type": "string",  
        "description": "The content of the message as  
a string."  
        },  
        {
```

```
"type": "array",

"items": {

  "type": "object",

  "properties": {

    "type": {

      "type": "string",

      "description": "Type of the content provided"

    },

    "text": {

      "type": "string"

    },

    "image_url": {
```

```
"type": "object",

"properties": {

  "url": {

    "type": "string",

    "pattern": "^data:*",

    "description": "image uri with data (e.g.
    ...). HTTP URL will
    not be accepted"

  }

}

}
```

```
"type": "object",

"properties": {

  "type": {

    "type": "string",

    "description": "Type of the content provided"
  },

  "text": {

    "type": "string"
  },

  "image_url": {

    "type": "object",

    "properties": {

      "url": {
```

```
"type": "string",

"pattern": "^data:*",

"description": "image uri with data (e.g.
...). HTTP URL will
not be accepted"

}

}

}

}

}

}

}

}],

"image": {
  "oneOf": [
    {
      "type": "array",
```

```
      "description": "An
array of integers that represent the image
data constrained to 8-bit unsigned integer
values. Deprecated, use image as a part of
messages now.",
      "items": {
        "type":
"number",
        "description":
"A value between 0 and 255"
      },
      {
        "type": "string",
        "format":
"binary",
        "description":
"Binary string representing the image
contents. Deprecated, use image as a part of
messages now."
      }
    ],
  },
  "functions": {
    "type": "array",
```

```
      "items": {
        "type": "object",
        "properties": {
          "name": {
            "type":
"string"

          },
          "code": {
            "type":
"string"

          }
        },
        "required": [
          "name",
          "code"
        ]
      },
      "tools": {
        "type": "array",
        "description": "A list of
tools available for the assistant to use.",
        "items": {
          "type": "object",
          "oneOf": [
```

```
        {  
            "properties":  
            {  
                "name": {  
  
                    "type": "string",  
  
                    "description": "The name of the tool. More  
descriptive the better."  
                },  
  
                "description": {  
  
                    "type": "string",  
  
                    "description": "A brief description of what  
the tool does."  
                },  
  
                "parameters": {  
  
                    "type": "object",  
  
                    "description": "Schema defining the parameters  
accepted by the tool.",
```



```
"properties": {  
  
  "type": {  
  
    "type": "string",  
  
    "description": "The type of the parameters  
object (usually 'object')."  
  },  
  
  "required": {  
  
    "type": "array",  
  
    "description": "List of required parameter  
names.",  
  
    "items": {  
  
      "type": "string"  
    }  
  },  
}
```

```
"properties": {  
  
  "type": "object",  
  
  "description": "Definitions of each  
parameter.",  
  
  "additionalProperties": {  
  
    "type": "object",  
  
    "properties": {  
  
      "type": {  
  
        "type": "string",  
  
        "description": "The data type of the  
parameter."  
  
      },  
  
      "description": {
```

```
"type": "string",

"description": "A description of the expected
parameter."

}

},

"required": [

"type",

"description"

]

}

},

"required": [

"type",
```

```
"properties"
    ]
    }
},
"required": [
    "name",

"description",

"parameters"
    ]
    },
    {
        "properties":
    {
        "type": {

"type": "string",

"description": "Specifies the type of tool
(e.g., 'function')."

        },

"function": {
```

```
"type": "object",

"description": "Details of the function
tool.",

"properties": {

  "name": {

    "type": "string",

    "description": "The name of the function."
  },

  "description": {

    "type": "string",

    "description": "A brief description of what
the function does."
  },

  "parameters": {
```

```
"type": "object",

"description": "Schema defining the parameters
accepted by the function.",

"properties": {

  "type": {

    "type": "string",

    "description": "The type of the parameters
object (usually 'object')."

  },

  "required": {

    "type": "array",

    "description": "List of required parameter
names.",

    "items": {
```

```
"type": "string"
```

```
}
```

```
},
```

```
"properties": {
```

```
"type": "object",
```

```
"description": "Definitions of each  
parameter.",
```

```
"additionalProperties": {
```

```
"type": "object",
```

```
"properties": {
```

```
"type": {
```

```
"type": "string",
```

```
"description": "The data type of the
```

```
parameter."  
  
},  
  
"description": {  
  
  "type": "string",  
  
  "description": "A description of the expected  
parameter."  
  
}  
  
},  
  
"required": [  
  
  "type",  
  
  "description"  
  
]  
  
}
```



```
}

},

"required": [

"type",

"properties"

]

}

},

"required": [

"name",

"description",

"parameters"

]

}

},
```

```
        "required": [  
            "type",  
            "function"  
        ],  
    },  
]  
}  
},  
"stream": {  
    "type": "boolean",  
    "default": false,  
    "description": "If true,  
the response will be streamed back  
incrementally."  
},  
"max_tokens": {  
    "type": "integer",  
    "default": 256,  
    "description": "The  
maximum number of tokens to generate in the  
response."  
},  
"temperature": {  
    "type": "number",  
    "default": 0.6,
```

```
    "minimum": 0,  
    "maximum": 5,  
    "description": "Controls  
the randomness of the output; higher values  
produce more random results."
```

```
  },  
  "top_p": {  
    "type": "number",  
    "minimum": 0,  
    "maximum": 2,  
    "description": "Controls  
the creativity of the AI's responses by  
adjusting how many possible words it  
considers. Lower values make outputs more  
predictable; higher values allow for more  
varied and creative responses."
```

```
  },  
  "top_k": {  
    "type": "integer",  
    "minimum": 1,  
    "maximum": 50,  
    "description": "Limits the  
AI to choose from the top 'k' most probable  
words. Lower values make responses more
```

focused; higher values introduce more variety and potential surprises."

```
    },  
    "seed": {  
      "type": "integer",  
      "minimum": 1,  
      "maximum": 9999999999,  
      "description": "Random  
seed for reproducibility of the generation."  
    },  
    "repetition_penalty": {  
      "type": "number",  
      "minimum": 0,  
      "maximum": 2,  
      "description": "Penalty  
for repeated tokens; higher values discourage  
repetition."  
    },  
    "frequency_penalty": {  
      "type": "number",  
      "minimum": 0,  
      "maximum": 2,  
      "description": "Decreases  
the likelihood of the model repeating the same  
lines verbatim."
```

```
    },  
    "presence_penalty": {  
      "type": "number",  
      "minimum": 0,  
      "maximum": 2,  
      "description": "Increases  
the likelihood of the model introducing new  
topics."  
    }  
  },  
  "required": [  
    "messages"  
  ]  
}  
]  
}
```

Was this helpful?



[Previous](#)

[← Dashboard](#)

[Next](#)

[Workers Bindings →](#)

Resources

[API](#)

[New to Cloudflare?](#)

[Directory](#)

[Sponsorships](#)

[Open Source](#)

Company

[cloudflare.com](#)

[Our team](#)

[Careers](#)

Support

[Help Center](#)

[System Status](#)

[Compliance](#)

[GDPR](#)

Tools

[Cloudflare Radar](#)

[Speed Test](#)

[Is BGP Safe Yet?](#)

[RPKI Toolkit](#)

[Certificate Transparency](#)

Community

 [X](#)

 [Discord](#)

 [YouTube](#)



© 2025 Cloudflare, Inc. • [Privacy Policy](#) • [Terms of Use](#)

• [Report Security Issues](#) • [Trademark](#) •  [Cookie Preferences](#)