



L

llava-1.5-7b-hf

Beta

Image-to-Text • llava-hf

@cf/llava-hf/llava-1.5-7b-hf

LLaVA is an open-source chatbot trained by fine-tuning LLaMA/Vicuna on GPT-generated multimodal instruction-following data. It is an auto-regressive language model, based on the transformer architecture.

Model Info

Beta

Yes

Usage

✓ Workers - TypeScript

TypeScript

```
export interface Env {  
  AI: Ai;  
}  
  
export default {  
  async fetch(request: Request, env: Env):  
    Promise<Response> {  
    const res = await  
    fetch("https://cataas.com/cat");  
    const blob = await res.arrayBuffer();  
    const input = {  
      image: [...new Uint8Array(blob)],  
      prompt: "Generate a caption for this  
image",  
      max_tokens: 512,  
    };  
    const response = await env.AI.run(  
      "@cf/llava-hf/llava-1.5-7b-hf",  
      input  
    );  
    return new  
    Response(JSON.stringify(response));  
  },  
}
```

```
} satisfies ExportedHandler<Env>;
```

Parameters

* indicates a required field

Input

- 0 **string**

Binary string representing the image contents.

- 1 **object**

- **image** **one of** required

- 0 **array**

An array of integers that represent the image data constrained to 8-bit unsigned integer values

- **items** **number**

A value between 0 and 255

- 1 **string**

Binary string representing the image contents.

- temperature **number**

Controls the randomness of the output; higher values produce more random results.

- prompt **string**

The input text prompt for the model to generate a response.

- raw **boolean**

If true, a chat template is not applied and you must adhere to the specific model's expected formatting.

- top_p **number**

Controls the creativity of the AI's responses by adjusting how many possible words it considers. Lower values make outputs more predictable; higher values allow for more varied and creative responses.

- top_k **number**

Limits the AI to choose from the top 'k' most probable words. Lower values make responses more focused; higher values introduce more variety and potential surprises.

- seed **number**

Random seed for reproducibility of the generation.

- repetition_penalty **number**

Penalty for repeated tokens; higher values discourage repetition.

- frequency_penalty **number**

Decreases the likelihood of the model repeating the same lines verbatim.

- presence_penalty **number**

Increases the likelihood of the model introducing new topics.

- max_tokens **integer** default 512

The maximum number of tokens to generate in the response.

Output

- description **string**

API Schemas

The following schemas are based on JSON Schema

Input

Output

```
{
  "oneOf": [
    {
      "type": "string",
      "format": "binary",
      "description": "Binary string
representing the image contents."
    },
    {
      "type": "object",
      "properties": {
        "image": {
          "oneOf": [
            {
              "type": "array",
              "description": "An
array of integers that represent the image
data constrained to 8-bit unsigned integer
values",
              "items": {
                "type":
"number",
                "description":
"A value between 0 and 255"
```

```
        }  
      },  
      {  
        "type": "string",  
        "format":  
"binary",  
        "description":  
"Binary string representing the image  
contents."  
      }  
    ]  
  },  
  "temperature": {  
    "type": "number",  
    "description": "Controls  
the randomness of the output; higher values  
produce more random results."  
  },  
  "prompt": {  
    "type": "string",  
    "description": "The input  
text prompt for the model to generate a  
response."  
  },  
  "raw": {
```

```
    "type": "boolean",  
    "default": false,  
    "description": "If true, a  
chat template is not applied and you must  
adhere to the specific model's expected  
formatting."
```

```
  },  
  "top_p": {  
    "type": "number",  
    "description": "Controls  
the creativity of the AI's responses by  
adjusting how many possible words it  
considers. Lower values make outputs more  
predictable; higher values allow for more  
varied and creative responses."
```

```
  },  
  "top_k": {  
    "type": "number",  
    "description": "Limits the  
AI to choose from the top 'k' most probable  
words. Lower values make responses more  
focused; higher values introduce more variety  
and potential surprises."
```

```
  },  
  "seed": {
```

```
      "type": "number",
      "description": "Random
seed for reproducibility of the generation."
    },
    "repetition_penalty": {
      "type": "number",
      "description": "Penalty
for repeated tokens; higher values discourage
repetition."
    },
    "frequency_penalty": {
      "type": "number",
      "description": "Decreases
the likelihood of the model repeating the same
lines verbatim."
    },
    "presence_penalty": {
      "type": "number",
      "description": "Increases
the likelihood of the model introducing new
topics."
    },
    "max_tokens": {
      "type": "integer",
      "default": 512,
```

```
"description": "The maximum number of tokens to generate in the response."
```

```
    }  
  },  
  "required": [  
    "image"  
  ]  
}  
]  
}
```

Was this helpful?



Previous

← Dashboard

Next

Workers Bindings →

Resources

[API](#)

Support

[Help Center](#)

[New to Cloudflare?](#)[System Status](#)[Directory](#)[Compliance](#)[Sponsorships](#)[GDPR](#)[Open Source](#)

Company

Tools

[cloudflare.com](#)[Cloudflare Radar](#)[Our team](#)[Speed Test](#)[Careers](#)[Is BGP Safe Yet?](#)[RPKI Toolkit](#)[Certificate Transparency](#)

Community

[!\[\]\(465772ce2fc0e39b7001e2580b915cc2_img.jpg\) X](#)[!\[\]\(dc0c40d45c42e86bc0669168926f812c_img.jpg\) Discord](#)[!\[\]\(b58c23cb5aab1cd63092eda333892cb9_img.jpg\) YouTube](#)[!\[\]\(488d36215f31304317ffb20d512ebb61_img.jpg\) GitHub](#)

© 2025 Cloudflare, Inc. • [Privacy Policy](#) • [Terms of Use](#)

• [Report Security Issues](#) • [Trademark](#) •  [Cookie Preferences](#)

