



qwen2.5-coder-32b-instruct

Text Generation • Qwen

@cf/qwen/qwen2.5-coder-32b-instruct

Qwen2.5-Coder is the latest series of Code-Specific Qwen large language models (formerly known as CodeQwen). As of now, Qwen2.5-Coder has covered six mainstream model sizes, 0.5, 1.5, 3, 7, 14, 32 billion parameters, to meet the needs of different developers. Qwen2.5-Coder brings the following improvements upon CodeQwen1.5:

Model Info

Context	32,768 tokens
Window	↗

Model Info

LoRA	Yes
Unit Pricing	\$0.66 per M input tokens, \$1.00 per M output tokens

Playground

Try out this model with Workers AI LLM Playground. It does not require any setup or authentication and an instant way to preview and test a model directly in the browser.

Launch the LLM Playground

Usage

✓ Worker - Streaming

TypeScript

```
export interface Env {  
  AI: Ai;  
}  
  
export default {  
  async fetch(request, env):  
    Promise<Response> {  
  
    const messages = [  
      { role: "system", content: "You are a  
friendly assistant" },  
      {  
        role: "user",  
        content: "What is the origin of the  
phrase Hello, World",  
      },  
    ];  
  
    const stream = await  
env.AI.run("@cf/qwen/qwen2.5-coder-32b-  
instruct", {  
      messages,  
      stream: true,  
    });
```

```
return new Response(stream, {  
  headers: { "content-type":  
"text/event-stream" },  
  });  
},  
} satisfies ExportedHandler<Env>;
```

✓ Worker

TypeScript

```
export interface Env {  
  AI: Ai;  
}  
  
export default {  
  async fetch(request, env):  
    Promise<Response> {  
  
    const messages = [  
      { role: "system", content: "You are a  
friendly assistant" },  
      {  
        role: "user",  
        content: "What is the origin of the  
phrase Hello, World",  
      },  
    ];  
  
    const response = await  
env.AI.run("@cf/qwen/qwen2.5-coder-32b-  
instruct", { messages });  
  
    return Response.json(response);  
  },  
} satisfies ExportedHandler<Env>;
```

▼ Python

Python

```
import os
import requests

ACCOUNT_ID = "your-account-id"
AUTH_TOKEN = os.environ.get("CLOUDFLARE_AUTH_TOKEN")

prompt = "Tell me all about PEP-8"
response = requests.post(
    f"https://api.cloudflare.com/client/v4/accounts/{ACCOUNT_ID}/workers-ai/qwen2.5-coder-32b-instruct",
    headers={"Authorization": f"Bearer {AUTH_TOKEN}"},
    json={
        "messages": [
            {"role": "system", "content": "You are a helpful assistant."},
            {"role": "user", "content": prompt}
        ]
    }
)
result = response.json()
print(result)
```

▼ curl

```
curl
```

```
https://api.cloudflare.com/client/v4/accounts/  
coder-32b-instruct \  
  -X POST \  
  -H "Authorization: Bearer $CLOUDFLARE_AUTH_TOKEN" \  
  -d '{ "messages": [{ "role": "system", "content": "You are a helpful assistant." },  
    "user", "content": "Why is pizza so good" } ] }'
```

① OpenAI compatible endpoints

Workers AI also supports OpenAI compatible API endpoints for `/v1/chat/completions` and `/v1/embeddings`. For more details, refer to [Configurations](#).

Parameters

* indicates a required field

Input

• **object**

- **prompt** **string** required min 1

The input text prompt for the model to generate a response.

- **lora** **string**

Name of the LoRA (Low-Rank Adaptation) model to fine-tune the base model.

- **response_format** **object**

- **type** **string**

- **json_schema**

- **raw** **boolean**

If true, a chat template is not applied and you must adhere to the specific model's expected formatting.

- **stream** **boolean**

If true, the response will be streamed back incrementally using SSE, Server Sent Events.

- `max_tokens` **integer** default 256

The maximum number of tokens to generate in the response.

- `temperature` **number** default 0.6 min 0 max 5

Controls the randomness of the output; higher values produce more random results.

- `top_p` **number** min 0 max 2

Adjusts the creativity of the AI's responses by controlling how many possible words it considers. Lower values make outputs more predictable; higher values allow for more varied and creative responses.

- `top_k` **integer** min 1 max 50

Limits the AI to choose from the top 'k' most probable words. Lower values make responses more focused; higher values introduce more variety and potential surprises.

- **seed** **integer** min 1 max 9999999999

Random seed for reproducibility of the generation.

- **repetition_penalty** **number** min 0 max 2

Penalty for repeated tokens; higher values discourage repetition.

- **frequency_penalty** **number** min 0 max 2

Decreases the likelihood of the model repeating the same lines verbatim.

- **presence_penalty** **number** min 0 max 2

Increases the likelihood of the model introducing new topics.

- ▶ **1** **object**

- **messages** **array** required

An array of message objects representing the conversation history.

- **items** **object**

- **role** **string** required

The role of the message sender (e.g., 'user', 'assistant', 'system', 'tool').

- **content** **string** required

The content of the message as a string.

- **functions** **array**

- **items** **object**

- **name** **string** required

- **code** **string** required

- **tools** **array**

A list of tools available for the assistant to use.

- **items** **one of**

- **0** **object**

- **name** **string** required

The name of the tool. More descriptive the better.

- **description** **string** required

A brief description of what the tool does.

- **parameters** **object** required

Schema defining the parameters accepted by the tool.

- **type** **string** required

The type of the parameters object (usually 'object').

- **required** **array**

List of required parameter names.

- **items** **string**

- **properties** **object** required

Definitions of each parameter.

- **additionalProperties** **object**

- **type** **string** required

The data type of the parameter.

- **description** **string** required

A description of the expected parameter.

■ 1 **object**

- **type** **string** required

Specifies the type of tool (e.g., 'function').

- **function** **object** required

Details of the function tool.

- **name** **string** required

The name of the function.

- **description** **string** required

A brief description of what the function does.

- **parameters** **object** required

Schema defining the parameters accepted by the function.

- **type** **string** required

The type of the parameters object (usually 'object').

- **required** **array**

List of required parameter names.

- **items** **string**

- **properties** **object** required

Definitions of each parameter.

- **additionalProperties** **object**

- **type** **string** required

The data type of the parameter.

- **description** **string** required

A description of the expected parameter.

- **response_format** **object**

- **type** **string**

- **json_schema**

- **raw** **boolean**

If true, a chat template is not applied and you must adhere to the specific model's expected formatting.

- **stream** **boolean**

If true, the response will be streamed back incrementally using SSE, Server Sent Events.

- **max_tokens** **integer** default 256

The maximum number of tokens to generate in the response.

- **temperature** **number** default 0.6 min 0 max 5

Controls the randomness of the output; higher values produce more random results.

- **top_p** **number** min 0 max 2

Adjusts the creativity of the AI's responses by controlling how many possible words it considers. Lower values make outputs more predictable; higher values allow for more varied and creative responses.

- `top_k` **integer** min 1 max 50

Limits the AI to choose from the top 'k' most probable words. Lower values make responses more focused; higher values introduce more variety and potential surprises.

- `seed` **integer** min 1 max 9999999999

Random seed for reproducibility of the generation.

- `repetition_penalty` **number** min 0 max 2

Penalty for repeated tokens; higher values discourage repetition.

- `frequency_penalty` **number** min 0 max 2

Decreases the likelihood of the model repeating the same lines verbatim.

- `presence_penalty` **number** min 0 max 2

Increases the likelihood of the model introducing new topics.

Output

• **0** **object**

- response **string** required

The generated text response from the model

- usage **object**

Usage statistics for the inference request

- prompt_tokens **number** 0

Total number of tokens in input

- completion_tokens **number** 0

Total number of tokens in output

- total_tokens **number** 0

Total number of input and output tokens

- tool_calls **array**

An array of tool calls requests made during the response generation

- items **object**

- arguments **object**

The arguments passed to be passed to the tool call request

- name **string**

The name of the tool to be called

- ▶ 1 **string**

API Schemas

The following schemas are based on JSON Schema

Input

Output

```
{
  "type": "object",
  "oneOf": [
    {
      "title":
"Qwen2_5_Coder_32B_Instruct_Prompt",
      "properties": {
        "prompt": {
          "type": "string",
          "minLength": 1,
          "description": "The input
text prompt for the model to generate a
response."
        },
        "lora": {
          "type": "string",
          "description": "Name of
the LoRA (Low-Rank Adaptation) model to fine-
tune the base model."
        },
        "response_format": {
          "title": "JSON Mode",
          "type": "object",
          "properties": {
```

```
        "type": {
          "type": "string",
          "enum": [
            "json_object",
            "json_schema"
          ]
        },
        "json_schema": {}
      },
      "raw": {
        "type": "boolean",
        "default": false,
        "description": "If true, a
chat template is not applied and you must
adhere to the specific model's expected
formatting."
      },
      "stream": {
        "type": "boolean",
        "default": false,
        "description": "If true,
the response will be streamed back
incrementally using SSE, Server Sent Events."
      },
```



```
"max_tokens": {  
  "type": "integer",  
  "default": 256,  
  "description": "The  
maximum number of tokens to generate in the  
response."  
},  
"temperature": {  
  "type": "number",  
  "default": 0.6,  
  "minimum": 0,  
  "maximum": 5,  
  "description": "Controls  
the randomness of the output; higher values  
produce more random results."  
},  
"top_p": {  
  "type": "number",  
  "minimum": 0,  
  "maximum": 2,  
  "description": "Adjusts  
the creativity of the AI's responses by  
controlling how many possible words it  
considers. Lower values make outputs more
```

predictable; higher values allow for more varied and creative responses."

```
},
```

```
"top_k": {
```

```
  "type": "integer",
```

```
  "minimum": 1,
```

```
  "maximum": 50,
```

```
  "description": "Limits the
```

AI to choose from the top 'k' most probable

words. Lower values make responses more

focused; higher values introduce more variety

and potential surprises."

```
},
```

```
"seed": {
```

```
  "type": "integer",
```

```
  "minimum": 1,
```

```
  "maximum": 9999999999,
```

```
  "description": "Random
```

seed for reproducibility of the generation."

```
},
```

```
"repetition_penalty": {
```

```
  "type": "number",
```

```
  "minimum": 0,
```

```
  "maximum": 2,
```

```
        "description": "Penalty  
for repeated tokens; higher values discourage  
repetition."  
    },  
    "frequency_penalty": {  
        "type": "number",  
        "minimum": 0,  
        "maximum": 2,  
        "description": "Decreases  
the likelihood of the model repeating the same  
lines verbatim."  
    },  
    "presence_penalty": {  
        "type": "number",  
        "minimum": 0,  
        "maximum": 2,  
        "description": "Increases  
the likelihood of the model introducing new  
topics."  
    }  
},  
"required": [  
    "prompt"  
]  
},
```

```
{
  "title":
  "Qwen2_5_Coder_32B_Instruct_Messages",
  "properties": {
    "messages": {
      "type": "array",
      "description": "An array
of message objects representing the
conversation history.",
      "items": {
        "type": "object",
        "properties": {
          "role": {
            "type":
            "string",
            "description":
            "The role of the message sender (e.g., 'user',
'assistant', 'system', 'tool')."
          },
          "content": {
            "type":
            "string",
            "description":
            "The content of the message as a string."
          }
        }
      }
    }
  }
}
```

```
    },  
    "required": [  
        "role",  
        "content"  
    ]  
}  
},  
"functions": {  
    "type": "array",  
    "items": {  
        "type": "object",  
        "properties": {  
            "name": {  
                "type":  
  
"string"  
  
                },  
            "code": {  
                "type":  
  
"string"  
  
                },  
            "type":  
        }  
    },  
    "required": [  
        "name",  
        "code"  
    ]  
}
```

```
    }  
  },  
  "tools": {  
    "type": "array",  
    "description": "A list of  
tools available for the assistant to use.",  
    "items": {  
      "type": "object",  
      "oneOf": [  
        {  
          "properties":  
{  
            "name": {  
  
"type": "string",  
  
"description": "The name of the tool. More  
descriptive the better."  
          },  
  
"description": {  
  
"type": "string",  
  
"description": "A brief description of what
```

the tool does."

},

"parameters": {

"type": "object",

"description": "Schema defining the parameters
accepted by the tool.",

"properties": {

"type": {

"type": "string",

"description": "The type of the parameters
object (usually 'object')."

},

"required": {

"type": "array",

"description": "List of required parameter

```
names.",  
  
"items": {  
  
"type": "string"  
  
}  
  
},  
  
"properties": {  
  
"type": "object",  
  
"description": "Definitions of each  
parameter.",  
  
"additionalProperties": {  
  
"type": "object",  
  
"properties": {  
  
"type": {  
  
"type": "string",
```



```
"description": "The data type of the  
parameter."
```

```
},
```

```
"description": {
```

```
"type": "string",
```

```
"description": "A description of the expected  
parameter."
```

```
}
```

```
},
```

```
"required": [
```

```
"type",
```

```
"description"
```

```
]
```

```
}  
  
    },  
  
    "required": [  
  
    "type",  
  
    "properties"  
    ],  
  },  
  "required": [  
    "name",  
  
    "description",  
  
    "parameters"  
  ],  
},  
{  
  "properties":  
  
  "type": {
```

```
"type": "string",

"description": "Specifies the type of tool
(e.g., 'function')."

    },

"function": {

"type": "object",

"description": "Details of the function
tool.",

"properties": {

"name": {

"type": "string",

"description": "The name of the function."
    },

"description": {
```

```
"type": "string",

"description": "A brief description of what
the function does."

},

"parameters": {

"type": "object",

"description": "Schema defining the parameters
accepted by the function.",

"properties": {

"type": {

"type": "string",

"description": "The type of the parameters
object (usually 'object')."

},


```

```
"required": {  
  
  "type": "array",  
  
  "description": "List of required parameter  
names.",  
  
  "items": {  
  
    "type": "string"  
  
  }  
  
},  
  
"properties": {  
  
  "type": "object",  
  
  "description": "Definitions of each  
parameter.",  
  
  "additionalProperties": {
```

```
"type": "object",

"properties": {

  "type": {

    "type": "string",

    "description": "The data type of the
parameter."

  },

  "description": {

    "type": "string",

    "description": "A description of the expected
parameter."

  }

},
```

```
"required": [
```

```
"type",
```

```
"description"
```

```
]
```

```
}
```

```
}
```

```
},
```

```
"required": [
```

```
"type",
```

```
"properties"
```

```
]
```

```
}
```

```
},
```

} ,


```
        "enum": [  
            "json_object",  
            "json_schema"  
        ],  
    },  
    "json_schema": {}  
}  
,  
"raw": {  
    "type": "boolean",  
    "default": false,  
    "description": "If true, a  
chat template is not applied and you must  
adhere to the specific model's expected  
formatting."  
},  
"stream": {  
    "type": "boolean",  
    "default": false,  
    "description": "If true,  
the response will be streamed back  
incrementally using SSE, Server Sent Events."  
},  
"max_tokens": {  
    "type": "integer",
```

```
        "default": 256,  
        "description": "The  
maximum number of tokens to generate in the  
response."  
    },  
    "temperature": {  
        "type": "number",  
        "default": 0.6,  
        "minimum": 0,  
        "maximum": 5,  
        "description": "Controls  
the randomness of the output; higher values  
produce more random results."  
    },  
    "top_p": {  
        "type": "number",  
        "minimum": 0,  
        "maximum": 2,  
        "description": "Adjusts  
the creativity of the AI's responses by  
controlling how many possible words it  
considers. Lower values make outputs more  
predictable; higher values allow for more  
varied and creative responses."  
    },
```

```
"top_k": {  
  "type": "integer",  
  "minimum": 1,  
  "maximum": 50,  
  "description": "Limits the  
AI to choose from the top 'k' most probable  
words. Lower values make responses more  
focused; higher values introduce more variety  
and potential surprises."  
},  
"seed": {  
  "type": "integer",  
  "minimum": 1,  
  "maximum": 9999999999,  
  "description": "Random  
seed for reproducibility of the generation."  
},  
"repetition_penalty": {  
  "type": "number",  
  "minimum": 0,  
  "maximum": 2,  
  "description": "Penalty  
for repeated tokens; higher values discourage  
repetition."  
},
```

```
    "frequency_penalty": {
      "type": "number",
      "minimum": 0,
      "maximum": 2,
      "description": "Decreases
the likelihood of the model repeating the same
lines verbatim."
    },
    "presence_penalty": {
      "type": "number",
      "minimum": 0,
      "maximum": 2,
      "description": "Increases
the likelihood of the model introducing new
topics."
    }
  },
  "required": [
    "messages"
  ]
}
```

Was this helpful?



Previous

← Dashboard

Next

Workers Bindings →

Resources

[API](#)

[New to Cloudflare?](#)

[Directory](#)

[Sponsorships](#)

[Open Source](#)

Support

[Help Center](#)

[System Status](#)

[Compliance](#)

[GDPR](#)

Company

[cloudflare.com](#)

[Our team](#)

[Careers](#)

Tools

[Cloudflare Radar](#)

[Speed Test](#)

[Is BGP Safe Yet?](#)

[RPKI Toolkit](#)

Certificate Transparency

Community



© 2025 Cloudflare, Inc. • [Privacy Policy](#) • [Terms of Use](#)

• [Report Security Issues](#) • [Trademark](#) •  [Cookie Preferences](#)