

[Directory](#) > [R2](#) >[R2 Data Catalog](#) >[Connect to Iceberg engines](#) >[Apache Trino](#)

Copy page



# Apache Trino

Below is an example of using [Apache Trino](#) to connect to R2 Data Catalog. For more information on connecting to R2 Data Catalog with Trino, refer to [Trino documentation](#).

## Prerequisites

- Sign up for a [Cloudflare account](#).
- [Create an R2 bucket](#) and [enable the data catalog](#).
- [Create an R2 API token, key, and secret](#) with both [R2 and data catalog permissions](#).

- Install [Docker](#) to run the Trino container.

# Setup

Create a local directory for the catalog configuration and change directories to it

```
mkdir -p trino-catalog && cd trino-catalog/
```

Create a configuration file called `r2.properties` for your R2 Data Catalog connection:

```
# r2.properties

connector.name=iceberg

# R2 Configuration
fs.native-s3.enabled=true
s3.region=auto
s3.aws-access-key=<Your R2 access key>
s3.aws-secret-key=<Your R2 secret>
s3.endpoint=<Your R2 endpoint>
s3.path-style-access=true

# R2 Data Catalog Configuration
iceberg.catalog.type=rest
iceberg.rest-catalog.uri=<Your R2 Data Catalog
URI>
iceberg.rest-catalog.warehouse=<Your R2 Data
Catalog warehouse>
iceberg.rest-catalog.security=OAUTH2
iceberg.rest-catalog.oauth2.token=<Your R2
authentication token>
```

## Example usage

1

## Start Trino with the R2 catalog configuration:

```
# Create a local directory for the  
catalog configuration  
mkdir -p trino-catalog  
  
# Place your r2.properties file in the  
catalog directory  
cp r2.properties trino-catalog/  
  
# Run Trino with the catalog  
configuration  
docker run -d \  
    --name trino-r2 \  
    -p 8080:8080 \  
    -v $(pwd)/trino-  
catalog:/etc/trino/catalog \  
    trinodb/trino:latest
```

2

## Connect to Trino and query your R2 Data Catalog:



```
# Connect to the Trino CLI
```

```
docker exec -it trino-r2 trino
```

3

In the Trino CLI, run the following commands:

*-- Show all schemas in the R2 catalog*

```
SHOW SCHEMAS IN r2;
```

*-- Show all schemas in the R2 catalog*

```
CREATE SCHEMA r2.example_schema
```

*-- Create a table with some values in it*

```
CREATE TABLE
```

```
r2.example_schema.yearly_clicks (
```

```
    year,
```

```
    clicks
```

```
)
```

```
WITH (
```

```
    partitioning = ARRAY['year']
```

```
)
```

```
AS VALUES
```

```
    (2021, 10000),
```

```
    (2022, 20000);
```

*-- Show tables in a specific schema*

```
SHOW TABLES IN r2.example_schema;
```

*-- Query your Iceberg table*

```
SELECT * FROM  
r2.example_schema.yearly_clicks;
```

Was this helpful?



Previous

← About compaction

Next

DuckDB →

Last updated: Sep 5, 2025

## Resources

[API](#)

[New to Cloudflare?](#)

[Directory](#)

[Sponsorships](#)

[Open Source](#)

## Support

[Help Center](#)

[System Status](#)

[Compliance](#)

[GDPR](#)

## Company

## Tools

[cloudflare.com](#)

[Cloudflare Radar](#)

[Our team](#)

[Speed Test](#)

[Careers](#)

[Is BGP Safe Yet?](#)

[RPKI Toolkit](#)

[Certificate Transparency](#)

## Community

 [X](#)

 [Discord](#)

 [YouTube](#)

 [GitHub](#)

© 2025 Cloudflare, Inc. • [Privacy Policy](#) • [Terms of Use](#)

• [Report Security Issues](#) • [Trademark](#) •  [Cookie Preferences](#)